

Assignment 3

SOFE4790/CSCI4640 Distributed Systems

Total mark: 45


Due: November 17

1. A significant concern in distributed systems is conflicts that might arise from unsynchronized requests to a single resource. A very simple example of this concern can be demonstrated by a class that increments a counter through a method. If this class is instantiated by another class that supports multi-threading, and increments the counter N times within its `run()` method one can demonstrate that when multiple threads of this class are created there are situations when the value of the total count is not what is expected. For example, let's say that the multi-threaded counter object increments count 10 times and 3 threads of this object type are launched. The expected value of count is 30 (10×3) but this will not be the case if the 3 threads are not synchronized.

For this question create a `Counter` class with 2 methods `increaseCount()` and `getCount()` that correspondingly increase the value of a counter by 1 and gets the value of the counter. Now define a `CountingThread` class that instantiates a `Counter` object and implements the `Runnable` interface (i.e. supports multi-threading). In the `run()` method of `CountingThread` call `increaseCount()` several times. In the main method of `CountingThread` instantiate 3 threads of `CountingThread`. When these threads end get the value of the counter using `getCount()` and print out the result. Note: you might have to put the `Counter` object to sleep for a few milliseconds in the `increaseCount()` method to get significant synchronization issues. Code this and show that you have issues when you do not use method-level synchronization. Add the method-level synchronization to the code so that the expected counter sum is achieved. [10]

2. A file server uses caching, and achieves a hit rate of 80%. File operations in the server cost 5 ms of CPU time when the server finds the requested block in the cache, and take an additional 15 ms of disk I/O time otherwise. Explaining any assumptions you make, estimate the server's throughput capacity (average requests/sec) if it is:
i) single-threaded;
ii) two-threaded, running on a single processor;
iii) two-threaded, running on a two-processor computer. [8]
3. A clock is reading 10:27:54.0 (hr:min:sec) when it is discovered to be 4 seconds fast. Explain why it is undesirable to set it back to the right time at that point and show (numerically) how it should be adjusted so as to be correct after 8 seconds has elapsed. [5]

4.



The figure above shows events occurring for each of two processes, p_1 and p_2 . Arrows between processes denote message transmission. Draw and label the lattice of consistent states (p_1 state, p_2 state), beginning with the initial state $(0,0)$. [15]

5. In a certain system, each process typically uses a critical section many times before another process requires it. Explain why Ricart and Agrawala's multicast-based mutual exclusion algorithm is inefficient for this case, and describe how to improve its performance. Does your adaptation satisfy liveness condition ME2? [7]